Abhishek Tiwari

Senior Data Engineer

Mumbai, Maharashtra, India | +91 8960077051 | abhi1931480@gmail.com | LinkedIn | GitHub

PROFESSIONAL SUMMARY

Senior Data Engineer with 5 years of experience designing scalable data platforms, distributed ETL pipelines, and cloud-based Lakehouse architectures. Expertise in AWS data services, Apache Spark, and Kafka. Proven record of reducing pipeline latency, cutting compute costs, and delivering curated data products at enterprise scale across fintech and retail domains. Experienced in mentoring engineers and collaborating across product, data science, and BI teams.

PROFESSIONAL EXPERIENCE

JP Morgan Chase & Co. — Senior Data Engineer

Mumbai | Feb 2026 – Present

- Architected and maintained enterprise-scale data lake (Raw → Trusted → Refined) for Credit Card Rewards and Airport Lounge Access analytics platform, serving 3+ downstream BI and ML teams.
- Built serverless event-driven ingestion workflows using AWS Lambda + AWS Glue + PySpark, processing millions of credit card transactions daily with end-to-end data lineage.
- Optimized Spark partition strategies and transformation logic, improving processing throughput by ~35% and reducing job runtime on large financial datasets.
- Engineered refined datasets in Snowflake consumed by business intelligence dashboards and ML models, reducing analyst query time by 40%.
- Automated CI/CD pipelines using Jenkins + Git for version-controlled deployment of ETL jobs, cutting manual deployment effort by 60%.
- Containerized transformation workloads via Kubernetes on Amazon EKS, enabling horizontal scaling to handle peak transaction volumes.
- Implemented AWS CloudWatch monitoring and alerting, achieving 99.5% pipeline uptime and reducing mean time to resolution for data incidents.
- Mentored 2 junior data engineers on Spark best practices, data lake governance, and CI/CD workflows.

IBM India Pvt. Limited — Data Engineer

Mumbai | Jun 2025 – Jan 2026

- Led migration of legacy Oracle data warehouse to Apache Iceberg-based Lakehouse, improving query performance by 50% and reducing storage costs for 150M+ daily transaction records.
- Built high-throughput Kafka + PySpark streaming ingestion pipelines processing 150M+ banking records daily with sub-minute latency.
- Designed schema evolution framework for Iceberg datasets, eliminating breaking changes across 10+ downstream reporting pipelines.
- Implemented data quality validation and automated reconciliation frameworks, reducing data discrepancy incidents by 30% across regulatory reports.
- Collaborated with cross-functional teams to design optimized data models for regulatory reporting and operational analytics.

LTIMindtree (Larsen and Toubro group Company) — Data Engineer

Noida | Aug 2021 – Jun 2025

- Developed scalable distributed ETL pipelines for Adidas Global Lakehouse platform processing 100TB+ of retail data across sales, supply chain, and customer domains.
- Designed batch and streaming PySpark pipelines enabling near real-time analytics, reducing reporting latency from hours to under 15 minutes.
- Integrated multi-source enterprise datasets across AWS and Azure hybrid environments, unifying analytics for 5+ global business units.
- Optimized Spark transformations and SQL queries, improving query performance by 45% for downstream BI dashboards.
- Built GenAI-powered RAG document assistant enabling analysts to query enterprise data lake using natural language, reducing ad-hoc data requests by 25%.
- Participated in data platform architecture reviews and contributed to enterprise data lake governance strategy across 3 client programs.

TECHNICAL SKILLS

Programming: Python, SQL, Java, Bash
Big Data: Apache Spark (PySpark), Kafka, Hadoop, Hive, Airflow, Iceberg, Delta Lake
Cloud & Platforms: AWS (S3, Glue, Lambda, EKS, EC2, IAM, CloudWatch), Databricks, Snowflake
Data Engineering: ETL Pipelines, Data Lake Architecture, Data Modeling, Data Warehousing
DevOps & Tools: Git, Jenkins, Docker, Kubernetes, CI/CD, Jira, Confluence
Analytics & Visualization: Power BI, Pandas, NumPy, Matplotlib
Generative AI & ML: RAG, LLM Applications, Prompt Engineering, Vector Embeddings, Semantic Search, FastAPI, Scikit-learn, PyTorch

PROJECTS

- Data Harbour — Containerized Lakehouse platform integrating PySpark, Hive, Delta Lake, Airflow, PostgreSQL, and MinIO. GitHub Link
- Apache Airflow + Kafka Streaming Platform — Real-time streaming ingestion pipelines with workflow orchestration. GitHub Link
- RAG-Based AI Search Assistant — Enterprise document intelligence system using LLM-powered semantic search over large knowledge bases. GitHub Link

EDUCATION

Bachelor of Science (Computer Science) — Thakur College of Science & Commerce, Mumbai University | 2021 | CGPA: 8.23

CERTIFICATIONS

AWS Cloud Quest: Cloud Practitioner • Advanced SQL (HackerRank) • Databricks Generative AI Fundamentals • Apache Spark & PySpark • IBM Spark Fundamentals II • Cloudera Data Warehouse (DENG 156) • IBM Data Analysis Using Python